


# CS 168 Fall 2016 Section 4 – Distance Vector Routing



## Problem 1: Distance-Vector Routing

The nodes in the above network communicate with each other using Distance-Vector routing. Below are the initial routing tables for each node, and a table showing the costs for each of their neighboring links.

In the routing tables, each row represents a neighbor and a column represents a destination. Each cell entry is of the format (shortest known distance to dest, next hop).

**Node A**

<i>Nbr</i>	<i>Cost</i>	To From	A	B	C	D
A	0	A	0, A	2, B	1, C	7, D
B	2	B	-	0	-	-
C	1	C	-	-	0	-
D	7	D	-	-	-	0

**Node B**

<i>Nbr</i>	<i>Cost</i>	To From	A	B	E
A	2	A	0	-	-
B	0	B	2, A	0, B	5, E
E	5	E	-	-	0

**Node C**

<i>Nbr</i>	<i>Cost</i>	To From	A	C	D
A	1	A	0	-	-
C	0	C	1, A	0, C	2, D
D	2	D	-	-	0

**Node D**

<i>Nbr</i>	<i>Cost</i>	To From	A	C	D	E
A	7	A	0	-	-	-
C	2	C	-	0	-	-
D	0	D	7, A	2, C	0, D	3, E
E	3	E	-	-	-	0

**Node E**

<i>Nbr</i>	<i>Cost</i>	To From	B	D	E
B	5	B	0	-	-
D	3	D	-	0	-
E	0	E	5, B	3, D	0, E

The following questions indicate events that happen consecutively. You can assume that there are no other packet exchanges than the ones specified.

A. *C sends its update to A and D.*

A.1. What information is contained in C's update?

**C sends its distance vector (A: 1, D: 2)**

A.2. What do the routing tables for A and D look like after receiving C's update? (You may not need to fill in all columns)

Node A							Node D									
<i>Nbr</i>	<i>Cost</i>		To From	A	B	C	D	<i>Nbr</i>	<i>Cost</i>		To From	A	C	D	E	
A	0		A	0, A	2, B	1, C	3, C	A	7		A	0	-	-	-	
B	2		B	-	0	-	-	C	2		C	1	0	2	-	
C	1		C	1	-	0	2	D	0		D	3, C	2, C	0, D	3, E	
D	7		D	-	-	-	0	E	3		E	-	-	-	0	

A.3. Which nodes among A and D are expected to send routing updates after receiving C's update?

**Since both A and D updated their shortest paths, they will both send routing updates.**

B. *A sends its update to B, C, and D.*

B.1. What information is contained in A's update?

**(B: 2, C: 1, D: 3)**

B.2. What do the routing tables for B, C, and D look like after receiving A's update (You may not need to fill in all columns)?

Node B						Node C						Node D										
<i>Nbr</i>	<i>Cost</i>		A	B	C	D	E	<i>Nbr</i>	<i>Cost</i>		A	B	C	D		<i>Nbr</i>	<i>Cost</i>	A	B	C	D	E
A	2	A	0	2	1	3	-	A	1	A	0	2	1	3		A	7	0	2	1	3	-
B	0	B	2, A	0	3, A	5, A	5, E	C	0	C	1, A	3, A	0, C	2, D		C	2	1	-	0	2	-
E	5	E	-	-	-	-	0	D	2	D	-	-	-	0		D	0	3, C	9, A	2, C	0, D	3, E
																E	5	-	-	-	-	0

B.3. At this point, what route does D use to reach B? It knows that it can route to A via C with total distance 3 and that A can reach B with distance 2. Should it use this information to

optimize the route to B or should it wait for an update from C?

$D \rightarrow A \rightarrow B$

It should *not* use this information now, as it has no way of knowing whether C is aware of the route to B through A and if C will be able to forward a packet with destination B to A. D needs to wait until it explicitly receives a distance vector from C confirming this, before it can optimize its route to B.

B.4. Which nodes among B, C, and D are expected to send routing updates after receiving A's update?

B, C, and D will all send routing updates since they all updated their distance vectors.

C. **Skip this part in section, and use it as review for the exams.**

*D sends its update to A, C, and E.*

C.1. What information is contained in D's update?

(A: 3, B: 9, C: 2, E: 3)

C.2. What do the routing tables for A, C, and E look like after receiving D's update? (You may not need to fill in all columns)

Node A		Node C		Node E	
Nbr	Cost	Nbr	Cost	Nbr	Cost
A	0	A	0	A	0
B	2	B	1	B	5
C	1	C	0	C	3
D	7	D	2	D	3

C.3. Which nodes among A, C, and E are expected to send routing updates after receiving D's update?


A, C, and E all send updates.

D. Have the routing tables converged? Why or why not?

No. There are still distance vector updates (other than heartbeats) in the network.

## Problem 2: Count-To-Infinity Problem

Consider this simple topology:



- A. What values will be in the routing tables when the system has stabilized (after many rounds)?  
For this question, assume each node advertises all its real distances to all its neighbors.

Node A						Node B						Node C					
Nbr	Cost		A	B	C	Nbr	Cost		A	B	C	Nbr	Cost		A	B	C
A	0	A	0, A	1, B	2, B	A	1	A	0	1	2	B	1	B	1	0	1
B	1	B	1	0	1	B	0	B	1, A	0, B	1, C	C	0	C	2, B	1, B	0, C
		C				C	1	C	2	1	0						

- B. Now suppose the link from A to B goes down, such that A is no longer reachable:

- B.1. B notices the link outage and updates its routing table. What does B's updated routing table look like?

Node B					
Nbr	Cost		A	B	C
A	$\infty$	A	0	1	2
B	0	B	3, C	0, B	1, C
C	1	C	2	1	0

- B.2. According to its routing table, what is the cost of B's minimum-cost path to A?  
**3, by routing through C.**

- C. B sends an update to C. What is C's routing table after receiving the update?

Node C					
Nbr	Cost		A	B	C
B	1	B	3	0	1
C	0	C	4, B	1, B	0, C

- D. After updating its table, C sends an update to B. What is B's routing table after receiving the update?

<i>Nbr</i>	<i>Cost</i>
A	$\infty$
B	0
C	1

**Node B**

	A	B	C
A	0	1	2
B	5, C	0, B	1, C
C	4	1	0

- E. How many updates are exchanged before the tables converge?  
*Infinitely many.*

### Problem 3: Poison Reverse

One solution to the count-to-infinity problem is "poison-reverse": if you are currently routing through a neighbor, tell that neighbor that your path to the destination has infinite cost.

- A. Consider the network topology in Problem 2. Assuming that **poison reverse** was used when exchanging route information, what does B's routing table look like before the link from A to B goes down?

<i>Nbr</i>	<i>Cost</i>
A	1
B	0
C	1

Node B

	A	B	C
A	0	1	∞
B	1, A	0, B	1, C
C	∞	1	0

- B. *B detects the link outage and sends an update to C.*


B.1. What information is contained in B's update?

*B sends an update containing (A:  $\infty$ , C: 1).*

B.2. What does C's routing table look like after receiving the update?

		Node C			
<i>Nbr</i>	<i>Cost</i>		A	B	C
B	1	B	$\infty$	0	1
C	0	C	$\infty$	1, B	0, C

C. Now consider a more complex topology, with stabilized routing tables for A, B, and C:



**Node A**

<i>Nbr</i>	<i>Cost</i>		A	B	C	D
A	0	A	0, A	1, B	1, C	2, B
B	1	B	1	0	1	1
C	1	C	1	1	0	2

**Node B**

<i>Nbr</i>	<i>Cost</i>		A	B	C	D
A	1	A	0	1	1	$\infty$
B	0	B	1, A	0	1, C	1, D
C	1	C	1	1	0	$\infty$
D	1	D	$\infty$	1	$\infty$	0

**Node C**

<i>Nbr</i>	<i>Cost</i>		A	B	C	D
A	1	A	0	1	1	2
B	1	B	1	0	1	1
C	0	C	1, A	1, B	0, C	2, B

Suppose the link between B and D goes down. B notices this change and sends an update to A.

C.1. What is A's routing table after processing B's update?

**B sends an update with the only new distance being {D:  $\infty$ }.**

**Node A**

<i>Nbr</i>	<i>Cost</i>		A	B	C	D
A	0	A	0, A	1, B	1, C	3, C
B	1	B	1	0	1	$\infty$
C	1	C	1	1	0	2

C.2. A then sends an update back to B. What is B's routing table after processing A's update?


A tells B that A has a path to D of length 3.

Node B						
<i>Nbr</i>	<i>Cost</i>		A	B	C	D
A	1	A	0	1	1	3
B	0	B	1, A	0	1, C	4, A
C	1	C	1	1	0	∞
D	1	D	∞	1	∞	0

C.3. How might you avoid the count-to-infinity problem here altogether?

Use link-state instead [actual reason for the invention of link-state!], use path-vector, avoid topologies that are prone to this problem, route on DAGs, come up with some other awesome solution.

#### Problem 4: Split Poison?



A. Assume that the routers use **split horizon**. Say that E sends its initial update (B: 1, D: 1) to D. Assuming that D has received no other updates, what does D now tell E about D's path to B?

Nothing. Split Horizon means that we never tell a neighbor about paths that go through that neighbor. So in this case, D doesn't tell E about its path to B.

B. Assume that the routers use **poisoned reverse**. Furthermore, assume that the routing tables haven't converged, and D believes its shortest path to B is D-A-B (length 3). D sends this update to E. Now, E sends its first update (D: 1, B: 1) to D. After recomputing its routes, D sends an update to E. In this update, what is the advertised distance to B?

D will tell E that its distance to B is infinitely long, because D's new shortest route goes through E.

C. Now assume that the routers use **split horizon and poisoned reverse**. After the same scenario as in 2), what distance to B does D advertise to E?

D still tells E that D is infinitely far away from B. Split horizon and poisoned reverse are not mutually exclusive but actually typically used at the same time. In this case, even though split horizon says to not tell your neighbor about your paths through it, D must give E some update about its new path to B even though this route goes through E. Namely, if D doesn't update E, E will still believe D uses a route to B of length 3. We can avoid this by using poisoned reverse, which specifies that D should actively "lie" to E and say that its distance to B is  $\infty$ .

- D. Consider the simple topology (A-B-C) from problem 2. After the routing tables have converged, link A-B goes down. When B sends C an update containing (A:  $\infty$ ), is this an act of **poisoning a route** or **poisoned reverse**?

B is **poisoning a route**. Namely, it tells C that its distance is  $\infty$ , not because B's new path goes through C, but because B actually has no route now.

- E. **Poisoning a route** and **poisoned reverse** might sound similar, but actually we can think of one of them as being "honest" while the other one is "lying." Which one tells the truth, and which one tells a white lie to keep the network functioning?

**Poisoned reverse** encourages routers to tell a white lie. With poisoned reverse, we tell a neighbor that we have *no* path to a certain destination if our path goes through that neighbor. Since we actually do have a path, our message is not strictly true.

On the other hand, **poisoning a route** happens when a link goes down, and we actually lose our path to some destination. Thus, we're telling the truth when we advertise a distance of  $\infty$  to this destination (given that a infinitely long path is equivalent to no path).